

Alexander Alexandrov

Technische Universität Berlin
Database Systems and
Information Management Group
Building EN 7, Room EN 728
Einsteinufer 17, 10587 Berlin, Germany

Phone: (+49) 30 314 23599
Fax: (+49) 30 314 21601
Email: alexander.alexandrov@tu-berlin.de
Homepage: <http://aalexandrov.name/>

Personal

Born on May 24, 1895 in Sofia, Bulgaria.

Bulgarian Citizen.

Fluent in German and English.

Education

MSc in Computer Science, Technische Universität Berlin, 2011, summa cum laude.

BSc in Computer Science, Universität Mannheim, 2008.

Employment

Research Assistant, Technische Universität Berlin 2011–2016.

Student Assistant, Technische Universität Berlin 2009–2011.

Web Developer, bitExpert AG 2007–2009.

Publications

For a full list of my publications please refer to my DBLP record.

Journal Articles

Implicit Parallelism through Deep Language Embedding. 2016, SIGMOD Record (to appear).

The Stratosphere platform for big data analytics, 2014, *VLDB Journal*.

Proceedings

Implicit Parallelism through Deep Language Embedding. SIGMOD 2015.

Benchmarking in the Cloud: What It Should, Can, and Cannot Be, TPCTC 2012.

Professional Service

I have served as an external reviewer for various peer-reviewed conferences and journals, among which SIGMOD, VLDB, and TODS.

Teaching

I have co-organized the following Master courses at TU Berlin.

Database Technology, (WS 2011, WS 2012, WS 2013, WS 2014, WS 2015)

Implementation of Database Systems, (WS 2011, WS 2012, WS 2013, WS 2014, WS 2015)

Big Data Analytics Project, (SS 2014, WS 2014, SS 2015, WS 2015)

Open Source

I am the lead developer of the following open source projects.

Emma. <http://emma-language.org>

A DSL for data analysis in Scala hides low-level parallelism aspects from the programmer.

Technologies: *Scala, Scala Macros, Spark, Flink.*

Peel Framework. <http://peel-framework.org>

A framework that helps you to define, execute, analyze, and share experiments for distributed systems and algorithms.

Technologies: *Scala, Spring.*

Myriad Data Generator Toolkit. <http://myriad-toolkit.com>

A development toolkit for scalable data generators.

Technologies: *C++, Python, XML.*

Interests

I am interested in bridging the gap between the demands of modern Big Data analysis platforms and the need for high-level, declarative analytics languages. The approach I take is to introduce fundamental ideas and methods from database query optimization into language compilation pipelines through metaprogramming.

I am also interested in methods and techniques for scalable data generation and benchmarking of data analysis platforms.

Outside the area of computer science, my interests are mainly in cinema and musical folklore from the Balkans and the Middle East.

Last updated: January 17, 2016

<http://aalexandrov.name/assets/pdf/cv.pdf>